

Deep Learning: Shallow Supervised Learning

Lecture 01

Razvan C. Bunescu

School of Electrical Engineering and Computer Science

bunescu@ohio.edu


Supervised Learning

- **Task** = learn an (unkown) function $t : X \rightarrow T$ that maps input instances $\mathbf{x} \in X$ to output targets $t(\mathbf{x}) \in T$:
 - **Classification**:
 - The output $t(\mathbf{x}) \in T$ is one of a finite set of discrete categories.
 - **Regression**:
 - The output $t(\mathbf{x}) \in T$ is continuous, or has a continuous component.
- Target function $t(\mathbf{x})$ is known (only) through (noisy) set of training examples:
$$(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots (\mathbf{x}_n, t_n)$$


Supervised Learning

- **Task** = learn an (unknown) function $t : X \rightarrow T$ that maps input instances $\mathbf{x} \in X$ to output targets $t(\mathbf{x}) \in T$:
 - function t is known (only) through (noisy) set of training examples:
 - Training/Test data: $(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots (\mathbf{x}_n, t_n)$
- **Task** = build a function $h(\mathbf{x})$ such that:
 - h matches t well on the *training data*:
 - $\Rightarrow h$ is able to fit data that it has seen.
 - h also matches target t well on *test data*:
 - $\Rightarrow h$ is able to generalize to unseen data.

Regression: Curve Fitting




Regression: Curve Fitting



- **Training:** Build a function $h(\mathbf{x})$, based on training examples $(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_N, t_N)$


Regression: Curve Fitting




- **Testing:** for arbitrary (unseen) instance $\mathbf{x} \in X$, compute target output $h(\mathbf{x})$; want it to be close to $t(\mathbf{x})$.

Supervised Learning

Training




Testing



Parametric Approaches to Supervised Learning

- **Task** = build a function $h(\mathbf{x})$ such that:
 - h matches t well on the training data:
=> h is able to fit data that it has seen.
 - h also matches t well on test data:
=> h is able to generalize to unseen data.
- **Task** = choose h from a “nice” *class of functions* that depend on a vector of parameters \mathbf{w} :
 - $h(\mathbf{x}) \equiv h_{\mathbf{w}}(\mathbf{x}) \equiv h(\mathbf{w}, \mathbf{x})$
 - **what classes of functions are “nice”?**

Neurons



Soma is the central part of the neuron:

- *where the input signals are combined.*

Dendrites are cellular extensions:

- *where majority of the input occurs.*

Axon is a fine, long projection:

- *carries nerve signals to other neurons.*

Synapses are molecular structures between axon terminals and other neurons:

- *where the communication takes place.*

Neuron Models

<https://www.research.ibm.com/software/IBMResearch/multimedia/IJCNN2013.neuron-model.pdf>

Year	Model Name	Reference
1907	Integrate and fire	[13]
1943	McCulloch and Pitts	[11]
1952	Hodgkin-Huxley	[12]
1958	Perceptron	[14]
1961	Fitzhugh-Nagumo	[15]
1965	Leaky integrate-and-fire	[16]
1981	Morris-Lecar	[17]
1986	Quadratic integrate-and-fire	[18]
1989	Hindmarsh-Rose	[19]
1998	Time-varying integrate-and-fire model	[20]
1999	Wilson Polynomial	[21]
2000	Integrate-and-fire or burst	[22]
2001	Resonate-and-fire	[23]
2003	Izhikevich	[24]
2003	Exponential integrate-and-fire	[25]
2004	Generalized integrate-and-fire	[26]
2005	Adaptive exponential integrate-and-fire	[27]
2009	Mihalas-Neibur	[28]

Spiking/LIF Neuron Function

<http://ee.princeton.edu/research/prucnal/sites/default/files/06497478.pdf>

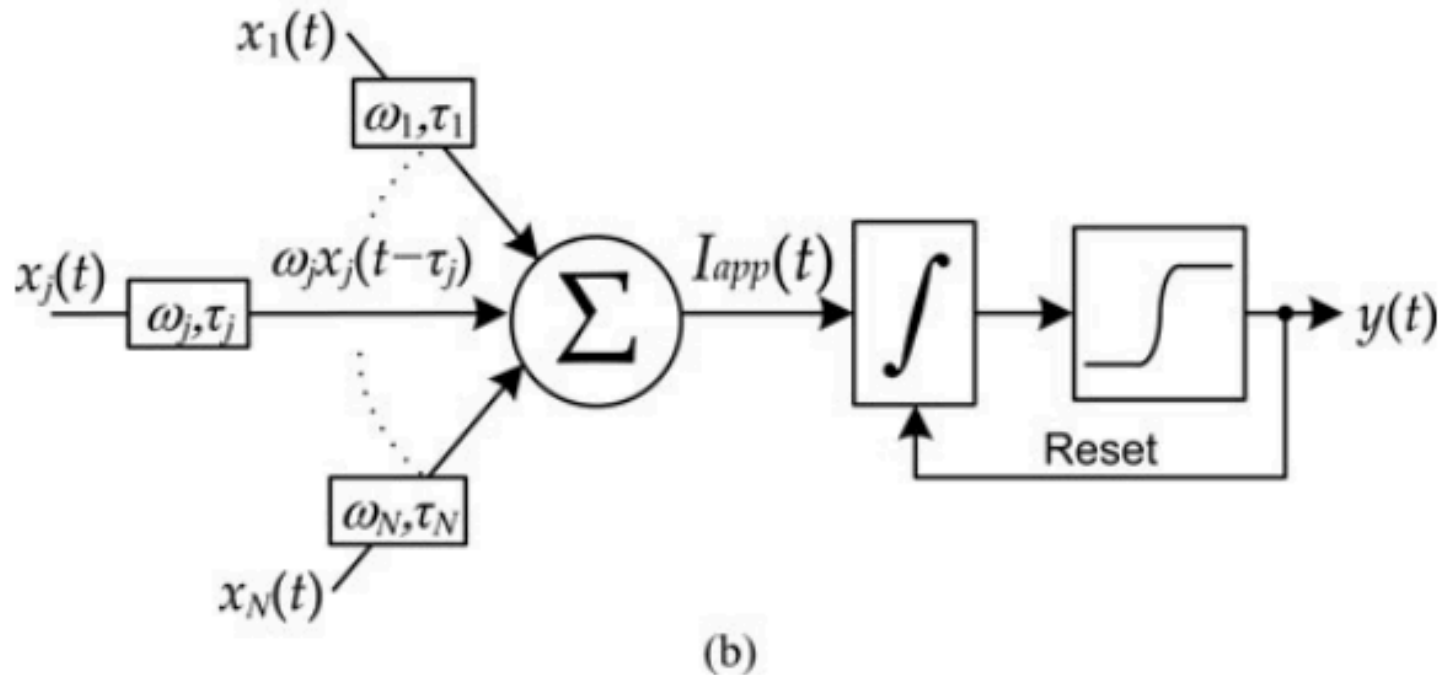



Fig. 2. (a) Illustration and (b) functional description of a leaky integrate-and-fire neuron. Weighted and delayed input signals are summed into the input current $I_{app}(t)$, which travel to the soma and perturb the internal state variable, the voltage V . Since V is hysteric, the soma performs integration and then applies a threshold to make a spike or no-spike decision. After a spike is released, the voltage V is reset to a value V_{reset} . The resulting spike is sent to other neurons in the network.

Neuron Models

<https://www.research.ibm.com/software/IBMResearch/multimedia/IJCNN2013.neuron-model.pdf>

Year	Model Name	Reference
1907	Integrate and fire	[13]
1943	McCulloch and Pitts	[11]
1952	Hodgkin-Huxley	[12]
1958	Perceptron	[14]
1961	Fitzhugh-Nagumo	[15]
1965	Leaky integrate-and-fire	[16]
1981	Morris-Lecar	[17]
1986	Quadratic integrate-and-fire	[18]
1989	Hindmarsh-Rose	[19]
1998	Time-varying integrate-and-fire model	[20]
1999	Wilson Polynomial	[21]
2000	Integrate-and-fire or burst	[22]
2001	Resonate-and-fire	[23]
2003	Izhikevich	[24]
2003	Exponential integrate-and-fire	[25]
2004	Generalized integrate-and-fire	[26]
2005	Adaptive exponential integrate-and-fire	[27]
2009	Mihalas-Neibur	[28]

McCulloch-Pitts Neuron Function



- Algebraic interpretation:
 - The output of the neuron is a **linear combination** of inputs from other neurons, **rescaled by** the synaptic **weights**.
 - weights w_i correspond to the synaptic weights (activating or inhibiting).
 - summation corresponds to combination of signals in the soma.
 - It is often transformed through an **activation / output function**.

Activation Functions

unit step $f(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$


Perceptron

logistic $f(z) = \frac{1}{1 + e^{-z}}$

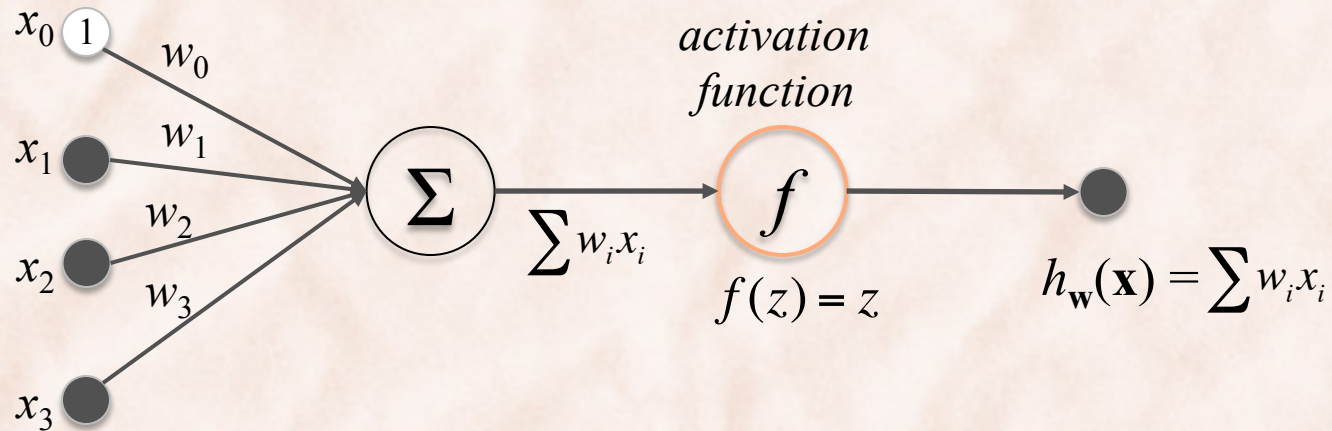
Logistic Regression

identity $f(z) = z$

Linear Regression



Linear Regression




- Polynomial curve fitting is Linear Regression:

$$\mathbf{x} = \phi(x) = [1, x, x^2, \dots, x^M]^T$$

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

Polynomial Curve Fitting



$$h(x) = h(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

↑ parameters

↑ features

Polynomial Curve Fitting

- Learning = finding the “right” parameters $\mathbf{w}^T = [w_0, w_1, \dots, w_M]$
 - Find \mathbf{w} that minimizes an *error / cost function* $E(\mathbf{w})$ which measures the misfit between $h(x_n, \mathbf{w})$ and t_n .
 - Expect that: $h(x_n, \mathbf{w})$ performing well on training examples $x_n \Rightarrow h(x, \mathbf{w})$ will perform well on arbitrary test examples $x \in X$.


Inductive Learning Hypothesis

- *Least Squares* error function:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{h(x_n, \mathbf{w}) - t_n\}^2$$

why squared?

Polynomial Curve Fitting



$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{h(x_n, \mathbf{w}) - t_n\}^2$$

- How do we find \mathbf{w}^* that minimizes $E(\mathbf{w})$?

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w})$$

Polynomial Curve Fitting

- *Least Square* solution is found by solving a set of $M + 1$ linear equations:

$$\sum_{j=0}^M A_{ij} w_j = T_i, \text{ where } A_{ij} = \sum_{n=1}^N x_n^{i+j}, \text{ and } T_i = \sum_{n=1}^N t_n x_n^i$$

- Prove it.

Gradient Descent (Batch)

- Want to minimize a function $f: R^n \rightarrow R$.
 - f is differentiable and convex.
 - compute gradient of f i.e. *direction of steepest increase*:

$$\nabla f(\mathbf{w}) = \left[\frac{df}{dw_1}(\mathbf{w}), \frac{df}{dw_2}(\mathbf{w}), \dots, \frac{df}{dw_k}(\mathbf{w}) \right]$$

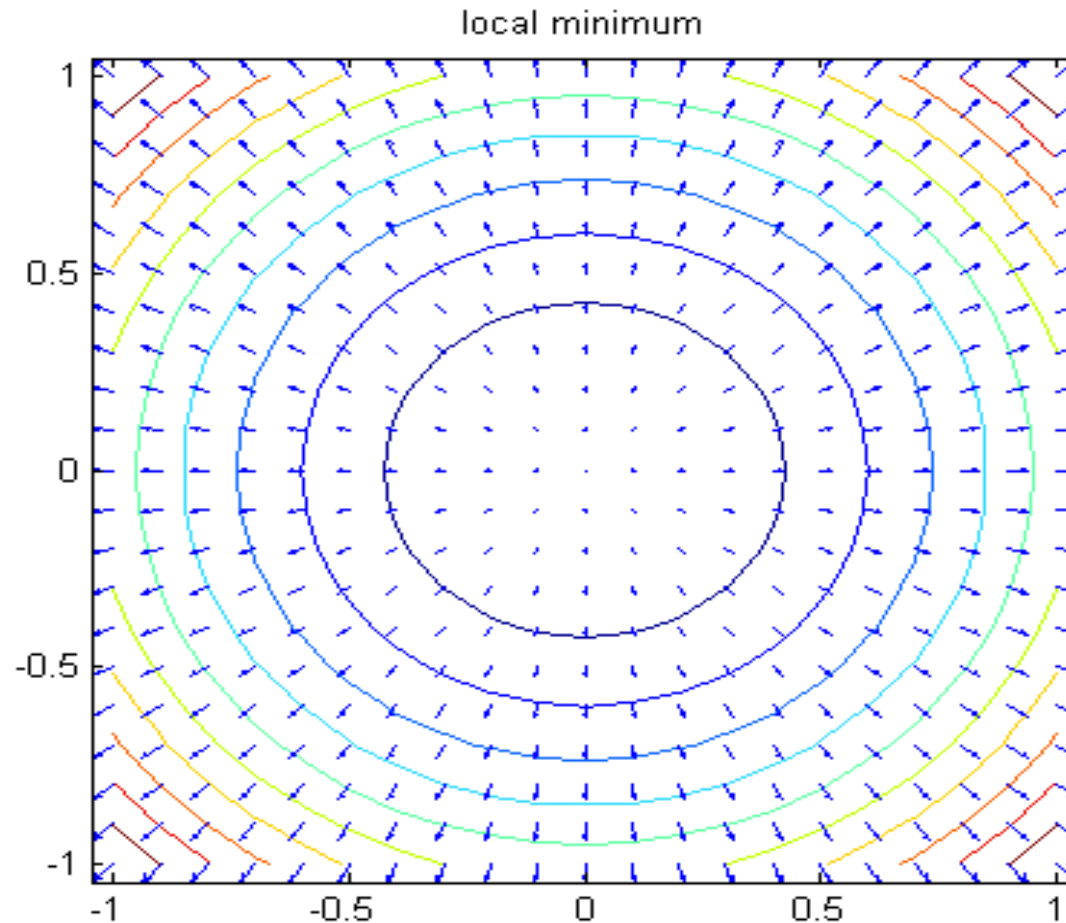
- choose a sequence of points $\mathbf{w}^1, \mathbf{w}^2, \dots$ and a learning rate η such that:

$$\mathbf{w}^{\tau+1} = \mathbf{w}^{\tau} - \eta \nabla f(\mathbf{w}^{\tau})$$

- Sum-of-squares error: $E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{\mathbf{w}^T \mathbf{x}_n - t_n\}^2$


Gradient Descent: Convex Objective

<http://www2.math.umd.edu/~jmr/241/gradients.html>




Gradient Descent: Non-Convex Objective

<http://www2.math.umd.edu/~jmr/241/gradients.html>




Gradient Descent: Saddle Points & Plateaus

<http://www2.math.umd.edu/~jmr/241/gradients.html>



Gradient Descent: Zig-Zagging Behavior



Stochastic Gradient Descent (Online)

- Decompose error function in sum of example errors:

$$E(\mathbf{w}) = \sum_{n=1}^N \boxed{\frac{1}{2} (\mathbf{w}^T \mathbf{x}_n - t_n)^2} = \sum_{n=1}^N E_n(\mathbf{w})$$

- Update parameters \mathbf{w} after each example, sequentially:


$$\begin{aligned} \mathbf{w}^{(\tau+1)} &= \mathbf{w}^{(\tau)} - \eta \nabla E_n(\mathbf{w}^{(\tau)}) \\ &= \mathbf{w}^{(\tau)} + \eta (t_n - \mathbf{w}^{(\tau)T} \mathbf{x}_n) \mathbf{x}_n \end{aligned}$$

=> the *least-mean-square* (LMS) algorithm.


Polynomial Curve Fitting

- **Generalization** = how well the parameterized $h(x, \mathbf{w}^*)$ performs on arbitrary (unseen) test instances $x \in X$.
- Generalization performance depends on the value of M .


0th Order Polynomial




1st Order Polynomial



3rd Order Polynomial



9th Order Polynomial



Polynomial Curve Fitting


- **Model Selection**: choosing the order M of the polynomial.
 - Best fit obtained with $M = 3$.
 - $M = 9$ obtains poor fit, even though it fits training examples perfectly:
 - But $M = 9$ polynomials subsume $M = 3$ polynomials!
- **Overfitting** \equiv good performance on training examples, poor performance on test examples.

Overfitting

- Measure fit using the Root-Mean-Square (RMS) error:

$$E_{RMS}(\mathbf{w}) = \sqrt{\frac{\sum_n (\mathbf{w}^T \mathbf{x}_n - t_n)^2}{N}}$$



- Use 100 random test examples, generated in the same way:



Over-fitting and Parameter Values

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Overfitting vs. Data Set Size



- More training data \Rightarrow less overfitting.
- What if we do not have more training data?
 - Use **regularization**.

Regularization

- **Parameter norm penalties** (term in the objective).
- Limit parameter norm (constraint).
- Dataset augmentation.
- Dropout.
- Ensembles.
- Semi-supervised learning.
- Early stopping.
- Noise robustness.
- Sparse representations.
- Adversarial training.


Regularization

- Penalize large parameter values:


$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{h(x_n, \mathbf{w}) - t_n\}^2 + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{regularizer}}$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w})$$


9th Order Polynomial with Regularization



9th Order Polynomial with Regularization



Training & Test error vs. $\ln \lambda$




How do we find the optimal value of λ ?


Model Selection

- Put aside an independent *validation set*.
- Select parameters giving best performance on validation set.

$$\ln \lambda \in \{-40, -35, -30, -25, -20, -15\}$$



Neuron Function



- Algebraic interpretation:
 - The output of the neuron is a **linear combination** of inputs from other neurons, **rescaled by** the synaptic **weights**.
 - weights w_i correspond to the synaptic weights (activating or inhibiting).
 - summation corresponds to combination of signals in the soma.
 - It is often transformed through a monotonic **activation function**.

Activation Functions

unit step $f(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$


Perceptron

logistic $f(z) = \frac{1}{1 + e^{-z}}$


Logistic Regression

identity $f(z) = z$

Linear Regression



Perceptron



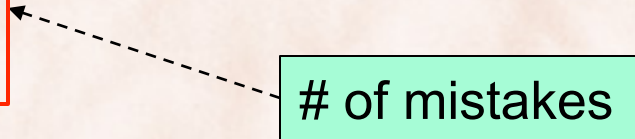
- Assume classes $T = \{c_1, c_2\} = \{1, 0\}$.
- Training set is $(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n)$.

$$\mathbf{x} = [1, x_1, x_2, \dots, x_k]^T$$

$$h(\mathbf{x}) = \text{step}(\mathbf{w}^T \mathbf{x})$$

Perceptron Learning

- Learning = finding the “right” parameters $\mathbf{w}^T = [w_0, w_1, \dots, w_k]$
 - Find \mathbf{w} that minimizes an *error function* $E(\mathbf{w})$ which measures the misfit between $h(\mathbf{x}_n, \mathbf{w})$ and t_n .
 - Expect that $h(\mathbf{x}, \mathbf{w})$ performing well on training examples $x_n \Rightarrow h(x, \mathbf{w})$ will perform well on arbitrary test examples $\mathbf{x} \in X$.
- **Least Squares** error function?

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{h(\mathbf{x}_n, \mathbf{w}) - t_n\}^2$$


of mistakes

Least Squares vs. Perceptron Criterion

- **Least Squares** \Rightarrow cost is # of misclassified patterns:
 - Piecewise constant function of \mathbf{w} with discontinuities.
 - Cannot find closed form solution for \mathbf{w} that minimizes cost.
 - Cannot use gradient methods (gradient zero almost everywhere).
- **Perceptron Criterion:**
 - Want $\mathbf{w}^T \mathbf{x}_n > 0$ for $t_n = 1$, and $\mathbf{w}^T \mathbf{x}_n < 0$ for $t_n = 0$.
 - \Rightarrow would like to have $\mathbf{w}^T \mathbf{x}_n (2t_n - 1) > 0$ for all patterns
 - \Rightarrow want to minimize $-\mathbf{w}^T \mathbf{x}_n (2t_n - 1)$ for all misclassified patterns M .

$$\Rightarrow \text{minimize } E_P(\mathbf{w}) = - \sum_{n \in M} \mathbf{w}^T \mathbf{x}_n (2t_n - 1)$$

Stochastic Gradient Descent

- Update parameters \mathbf{w} sequentially after each mistake:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_P(\mathbf{w}^{(\tau)}, \mathbf{x}_n)$$

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \eta \mathbf{x}_n (2t_n - 1)$$

- The magnitude of \mathbf{w} is inconsequential \Rightarrow set $\eta = 1$.

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \mathbf{x}_n (2t_n - 1)$$

The Perceptron Algorithm: Two Classes

1. **initialize** parameters $\mathbf{w} = 0$
2. **for** $i = 1 \dots n$
3. $h_i = \text{step}(\mathbf{w}^T \mathbf{x}_i)$
4. **if** $y_i \neq t_i$ **then**
5. $\mathbf{w} = \mathbf{w} + \mathbf{x}_i(2t_i - 1)$

Repeat:

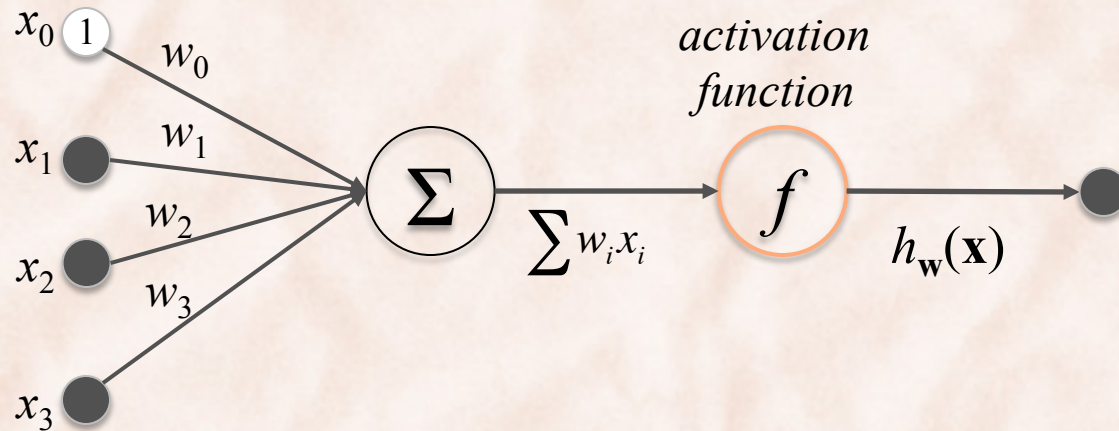
- a) until convergence.
- b) for a number of epochs E .

Theorem [[Rosenblatt, 1962](#)]:

If the training dataset is linearly separable, the perceptron learning algorithm is guaranteed to find a solution in a finite number of steps.

- see Theorem 1 (Block, Novikoff) in [[Freund & Schapire, 1999](#)].

Neuron Function



- Algebraic interpretation:
 - The output of the neuron is a **linear combination** of inputs from other neurons, **rescaled by** the synaptic **weights**.
 - weights w_i correspond to the synaptic weights (activating or inhibiting).
 - summation corresponds to combination of signals in the soma.
 - It is often transformed through a monotonic **activation function**.

Activation Functions

unit step $f(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$


Perceptron

logistic $f(z) = \frac{1}{1 + e^{-z}}$


Logistic Regression

identity $f(z) = z$

Linear Regression



Logistic Regression



- Can be used for both classification and regression:
 - **Classification:** $T = \{C_1, C_2\} = \{1, 0\}$.
 - **Regression:** $T = [0, 1]$ (i.e. output needs to be normalized).
- Training set is $(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n)$.

$$\mathbf{x} = [1, x_1, x_2, \dots, x_k]^T$$

$$h(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$$

Logistic Regression for Binary Classification

- Model output can be interpreted as **posterior class probabilities**:

$$p(C_1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

$$p(C_2 | \mathbf{x}) = 1 - \sigma(\mathbf{w}^T \mathbf{x}) = \frac{\exp(-\mathbf{w}^T \mathbf{x})}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

- How do we train a logistic regression model?
 - What **error/cost function** to minimize?

Logistic Regression Learning

- Learning = finding the “right” parameters $\mathbf{w}^T = [w_0, w_1, \dots, w_k]$
 - Find \mathbf{w} that minimizes an *error function* $E(\mathbf{w})$ which measures the misfit between $h(\mathbf{x}_n, \mathbf{w})$ and t_n .
 - Expect that $h(\mathbf{x}, \mathbf{w})$ performing well on training examples $\mathbf{x}_n \Rightarrow h(\mathbf{x}, \mathbf{w})$ will perform well on arbitrary test examples $\mathbf{x} \in X$.
- **Least Squares** error function?

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{h(\mathbf{x}_n, \mathbf{w}) - t_n\}^2$$

- Differentiable \Rightarrow can use gradient descent ✓
- Non-convex \Rightarrow not guaranteed to find the global optimum ✗

Maximum Likelihood

Training set is $D = \{\langle \mathbf{x}_n, t_n \rangle \mid t_n \in \{0,1\}, n \in 1 \dots N\}$

Let $h_n = p(C_1 \mid \mathbf{x}_n) \Leftrightarrow h_n = p(t_n = 1 \mid \mathbf{x}_n) = \sigma(\mathbf{w}^T \mathbf{x}_n)$

Maximum Likelihood (ML) principle: find parameters that maximize the likelihood of the labels.

- The **likelihood function** is $p(\mathbf{t} \mid \mathbf{w}) = \prod_{n=1}^N h_n^{t_n} (1 - h_n)^{(1-t_n)}$
- The negative log-likelihood (cross entropy) **error function**:

$$E(\mathbf{w}) = -\ln p(\mathbf{t} \mid \mathbf{x}) = -\sum_{n=1}^N \{t_n \ln h_n + (1 - t_n) \ln(1 - h_n)\}$$

Maximum Likelihood Learning for Logistic Regression

- The **ML** solution is:

$$\mathbf{w}_{ML} = \arg \max_{\mathbf{w}} p(\mathbf{t} | \mathbf{w}) = \arg \min_{\mathbf{w}} E(\mathbf{w})$$

convex in \mathbf{w}

- **ML** solution is given by $\nabla E(\mathbf{w}) = 0$.
 - Cannot solve analytically => solve numerically with gradient based methods: (stochastic) gradient descent, conjugate gradient, L-BFGS, etc.
 - Gradient is (prove it):

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (h_n - t_n) \mathbf{x}_n^T$$

Regularized Logistic Regression

- Use a Gaussian prior over the parameters:

$$\mathbf{w} = [w_0, w_1, \dots, w_M]^T$$

$$p(\mathbf{w}) = N(\mathbf{0}, \alpha^{-1}\mathbf{I}) = \left(\frac{\alpha}{2\pi}\right)^{(M+1)/2} \exp\left\{-\frac{\alpha}{2}\mathbf{w}^T\mathbf{w}\right\}$$

- Bayes' Theorem:

$$p(\mathbf{w} | \mathbf{t}) = \frac{p(\mathbf{t} | \mathbf{w})p(\mathbf{w})}{p(\mathbf{t})} \propto p(\mathbf{t} | \mathbf{w})p(\mathbf{w})$$

- MAP solution:

$$\mathbf{w}_{MAP} = \arg \max_{\mathbf{w}} p(\mathbf{w} | \mathbf{t})$$

Regularized Logistic Regression

- MAP solution:

$$\begin{aligned}\mathbf{w}_{MAP} &= \arg \max_{\mathbf{w}} p(\mathbf{w} | \mathbf{t}) = \arg \max_{\mathbf{w}} p(\mathbf{t} | \mathbf{w}) p(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} -\ln p(\mathbf{t} | \mathbf{w}) p(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} -\ln p(\mathbf{t} | \mathbf{w}) - \ln p(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} E_D(\mathbf{w}) - \ln p(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} E_D(\mathbf{w}) + \frac{\alpha}{2} \mathbf{w}^T \mathbf{w} \quad = \arg \min_{\mathbf{w}} E_D(\mathbf{w}) + E_w(\mathbf{w})\end{aligned}$$

$$E_D(\mathbf{w}) = -\sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln(1 - y_n)\} \quad \text{-----} \rightarrow \text{data term}$$

$$E_w(\mathbf{w}) = \frac{\alpha}{2} \mathbf{w}^T \mathbf{w} \quad \text{-----} \rightarrow \text{regularization term}$$

Regularized Logistic Regression

- MAP solution:

$$\mathbf{w}_{MAP} = \arg \min_{\mathbf{w}} E_D(\mathbf{w}) + E_w(\mathbf{w})$$

still convex in \mathbf{w}

- ML solution is given by $\nabla E(\mathbf{w}) = 0$.

$$\nabla E(\mathbf{w}) = \nabla E_D(\mathbf{w}) + \nabla E_w(\mathbf{w}) = \sum_{n=1}^N (h_n - t_n) \mathbf{x}_n^T + \alpha \mathbf{w}^T$$

where $h_n = \sigma(\mathbf{w}^T \mathbf{x}_n)$

- Cannot solve analytically \Rightarrow solve numerically:
 - (stochastic) gradient descent [PRML 3.1.3], Newton Raphson iterative optimization [PRML 4.3.3], conjugate gradient, LBFGS.

Softmax Regression = Logistic Regression for Multiclass Classification

- Multiclass classification:

$$T = \{C_1, C_2, \dots, C_K\} = \{1, 2, \dots, K\}.$$

- Training set is $(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_n, t_n)$.

$$\mathbf{x} = [1, x_1, x_2, \dots, x_M]$$

$$t_1, t_2, \dots, t_n \in \{1, 2, \dots, K\}$$

- One weight vector per class [[PRML 4.3.4](#)]:

$$p(C_k | \mathbf{x}) = \frac{\exp(\mathbf{w}_k^T \mathbf{x})}{\sum_j \exp(\mathbf{w}_j^T \mathbf{x})}$$

Softmax Regression ($K \geq 2$)

- Inference:

$$\begin{aligned} C_* &= \arg \max_{C_k} p(C_k | \mathbf{x}) \\ &= \arg \max_{C_k} \frac{\exp(\mathbf{w}_k^T \mathbf{x})}{\sum_j \exp(\mathbf{w}_j^T \mathbf{x})} \end{aligned}$$

$Z(\mathbf{x})$ a normalization constant

$$\begin{aligned} &= \arg \max_{C_k} \exp(\mathbf{w}_k^T \mathbf{x}) \\ &= \arg \max_{C_k} \mathbf{w}_k^T \mathbf{x} \end{aligned}$$

- Training using:

- Maximum Likelihood (ML)
- Maximum A Posteriori (MAP) with a Gaussian prior on \mathbf{w} .

Softmax Regression

- The **negative log-likelihood** error function is:

$$E_D(\mathbf{w}) = -\frac{1}{N} \ln \prod_{n=1}^N p(t_n | \mathbf{x}_n)$$

$$= -\frac{1}{N} \sum_{n=1}^N \ln \frac{\exp(\mathbf{w}_{t_n}^T \mathbf{x}_n)}{Z(\mathbf{x}_n)}$$

$$= -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K \delta_k(t_n) \ln \frac{\exp(\mathbf{w}_k^T \mathbf{x}_n)}{Z(\mathbf{x}_n)}$$

convex in \mathbf{w}

where $\delta_t(x) = \begin{cases} 1 & x = t \\ 0 & x \neq t \end{cases}$ is the *Kronecker delta* function.

Softmax Regression

- The **ML** solution is:

$$\mathbf{w}_{ML} = \arg \min_{\mathbf{w}} E_D(\mathbf{w})$$

- The **gradient** is (prove it):

$$\begin{aligned}\nabla_{\mathbf{w}_k} E_D(\mathbf{w}) &= -\frac{1}{N} \sum_{n=1}^N (\delta_k(t_n) - p(C_k | \mathbf{x}_n)) \mathbf{x}_n \\ &= -\frac{1}{N} \sum_{n=1}^N \left(\delta_k(t_n) - \frac{\exp(\mathbf{w}_k^T \mathbf{x}_n)}{Z(\mathbf{x}_n)} \right) \mathbf{x}_n\end{aligned}$$

$$\nabla E_D(\mathbf{w}) = \left[\nabla_{\mathbf{w}_1}^T E_D(\mathbf{w}), \nabla_{\mathbf{w}_2}^T E_D(\mathbf{w}), \dots, \nabla_{\mathbf{w}_K}^T E_D(\mathbf{w}) \right]^T$$

Regularized Softmax Regression

- The new **cost** function is:

$$E(\mathbf{w}) = E_D(\mathbf{w}) + E_w(\mathbf{w})$$

$$= -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K \delta_k(t_n) \ln \frac{\exp(\mathbf{w}_k^T \mathbf{x}_n)}{Z(\mathbf{x}_n)} + \frac{\alpha}{2} \sum_{k=1}^K \mathbf{w}_k^T \mathbf{w}_k$$

- The new **gradient** is (prove it):

$$\nabla_{\mathbf{w}_k} E(\mathbf{w}) = -\frac{1}{N} \sum_{n=1}^N (\delta_k(t_n) - p(C_k | \mathbf{x}_n)) \mathbf{x}_n^T + \alpha \mathbf{w}_k^T$$

Softmax Regression

- **ML** solution is given by $\nabla E_D(\mathbf{w}) = 0$.
 - Cannot solve analytically.
 - Solve numerically, by plugging $[cost, gradient] = [E_D(\mathbf{w}), \nabla E_D(\mathbf{w})]$ values into general convex solvers:
 - L-BFGS
 - Newton methods
 - conjugate gradient
 - (stochastic / minibatch) gradient-based methods.
 - gradient descent (with / without momentum).
 - AdaGrad, AdaDelta
 - RMSProp
 - ADAM, ...

Implementation

- Need to compute [*cost*, *gradient*]:

- $cost = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K \delta_k(t_n) \ln p(C_k | \mathbf{x}_n) + \frac{\alpha}{2} \sum_{k=1}^K \mathbf{w}_k^T \mathbf{w}_k$

- $gradient_k = -\frac{1}{N} \sum_{n=1}^N (\delta_k(t_n) - p(C_k | \mathbf{x}_n)) \mathbf{x}_n^T + \alpha \mathbf{w}_k^T$

=> need to compute, for $k = 1, \dots, K$:

- $output \ p(C_k | \mathbf{x}_n) = \frac{\exp(\mathbf{w}_k^T \mathbf{x}_n)}{\sum_j \exp(\mathbf{w}_j^T \mathbf{x}_n)}$

Overflow when $\mathbf{w}_k^T \mathbf{x}_n$ are too large.

Implementation: Preventing Overflows

- Subtract from each product $\mathbf{w}_k^T \mathbf{x}_n$ the maximum product:

$$c = \max_{1 \leq k \leq K} \mathbf{w}_k^T \mathbf{x}_n$$

$$p(C_k | \mathbf{x}_n) = \frac{\exp(\mathbf{w}_k^T \mathbf{x}_n - c)}{\sum_j \exp(\mathbf{w}_j^T \mathbf{x}_n - c)}$$

Implementation: Gradient Checking

- Want to minimize $J(\theta)$, where θ is a scalar.

- Mathematical definition of derivative:

$$\frac{d}{d\theta} J(\theta) = \lim_{\varepsilon \rightarrow 0} \frac{J(\theta + \varepsilon) - J(\theta - \varepsilon)}{2\varepsilon}$$

- Numerical approximation of derivative:

$$\frac{d}{d\theta} J(\theta) \approx \frac{J(\theta + \varepsilon) - J(\theta - \varepsilon)}{2\varepsilon} \quad \text{where } \varepsilon = 0.0001$$

Implementation: Gradient Checking

- If θ is a vector of parameters θ_i ,
 - Compute numerical derivative with respect to each θ_i .
 - Aggregate all derivatives into numerical gradient $G_{\text{num}}(\theta)$.
- Compare numerical gradient $G_{\text{num}}(\theta)$ with implementation of gradient $G_{\text{imp}}(\theta)$:

$$\frac{\|G_{\text{num}}(\theta) - G_{\text{imp}}(\theta)\|}{\|G_{\text{num}}(\theta) + G_{\text{imp}}(\theta)\|} \leq 10^{-6}$$